



DGIST

# Deep Learning Seminar

## Chapter 7. Regularization for Deep Learning

- Part 2 -

Hyun-Lim YANG

Department of Information and Communication Engineering

DGIST

2017.08.09



# ***Contents***

# Chapter 7. Regularization for Deep Learning

## ● Part 1

- 7.1 Parameter Norm Penalties
- 7.2 Norm Penalties as Constrained Optimization
- 7.3 Regularization and Under-Constrained Problems
- 7.4 Dataset Augmentations
- 7.5 Noise Robustness
- 7.6 Semi-Supervised Learning
- 7.7 Multi-Task Learning
- 7.8 Early Stopping

## ● Part 2

- 7.9 Parameter Tying and Parameter Sharing
- 7.10 Sparse Representations
- 7.11 Bagging and Other Ensemble Methods
- 7.12 Dropout
- 7.13 Adversarial Training
- 7.14 Tangent Distance, Tangent Prop, and Manifold Tangent Classifier

# Chapter 7. Regularization for Deep Learning

## ● Part 1

- 7.1 Parameter Norm Penalties
- 7.2 Norm Penalties as Constrained Optimization
- 7.3 Regularization and Under-Constrained Problems
- 7.4 Dataset Augmentations
- 7.5 Noise Robustness
- 7.6 Semi-Supervised Learning
- 7.7 Multi-Task Learning
- 7.8 Early Stopping

## ● Part 2

- 7.9 Parameter Tying and Parameter Sharing
- 7.10 Sparse Representations
- 7.11 Bagging and Other Ensemble Methods
- 7.12 Dropout
- 7.13 Adversarial Training
- 7.14 Tangent Distance, Tangent Prop, and Manifold Tangent Classifier

# Chapter 7. Regularization for Deep Learning

## ● Part 1

- 7.1 Parameter Norm Penalties
- 7.2 Norm Penalties as Constrained Optimization
- 7.3 Regularization and Under-Constrained Problems
- 7.4 Dataset Augmentations
- 7.5 Noise Robustness
- 7.6 Semi-Supervised Learning
- 7.7 Multi-Task Learning
- 7.8 Early Stopping

## ● Part 2

- 7.9 Parameter Tying and Parameter Sharing
- ~~7.10 Sparse Representations~~ Discuss in Chapter 14. Autoencoder
- 7.11 Bagging and Other Ensemble Methods
- 7.12 Dropout
- 7.13 Adversarial Training
- ~~7.14 Tangent Distance, Tangent Prop, and Manifold Tangent Classifier~~ Chapter 14.

# ***Parameter Tying and Parameter Sharing***

# Parameter Tying

## ● Parameter dependency

- $L^2$  regularization (or weigh decay) penalizes model parameters for deviating from the fixed value of zero
- Sometimes we need other ways to **express prior knowledge** of parameters
- We may know from domain and model architecture that there should be some dependencies between model parameters

## ● The goal of parameter tying

- We want to express that certain parameters should be close to one another

# A scenario of parameter tying

- Two models performing the same classification task (with same set of classes) but with somewhat different input distributions
- Model  $A$  with parameter  $w^{(A)}$
- Model  $B$  with parameter  $w^{(B)}$
- The two models will map the input to two different, but related output:

$$\hat{y}^{(A)} = f(w^{(A)}, x)$$

$$\hat{y}^{(B)} = g(w^{(B)}, x')$$



# $L^2$ penalty for parameter tying

- If the tasks are similar enough (perhaps with similar input and output distributions) then we believe that the model parameters should be close to each other:

$$\forall i, w_i^{(A)} \approx w_i^{(B)}$$

- We can leverage this information via regularization
- Use a parameter norm penalty (other choices are also possible)

$$\Omega(w^{(A)}, w^{(B)}) = \|w^{(A)} - w^{(B)}\|_2^2$$

Regularized objective function

Penalty term

$$\tilde{J}(\theta; X, y) = J(\theta; X, y) + \alpha \Omega(\theta)$$

Original objective function

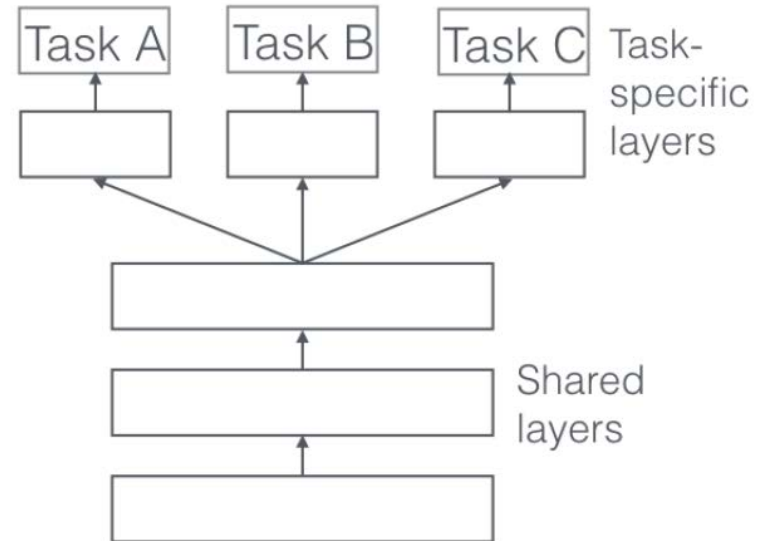
# Multi-Task Learning (MTL)

- Sharing the representation between related tasks
  - We can enable our model to **generalize better** on our original task
  - Another approach of bagging (with different cost functions)
- MTL is also known as:
  - Joint learning, Learning to learn, learning with auxiliary tasks
- Optimizing more than one loss function
- Improves generalization by leveraging the domain-specific information contained in the training data
  - Even if the problem optimizing one loss functions, there might be the chances to **improve performance by adding an auxiliary task upon the major task**
- Motivated by human learning process

# Two MTL Methods

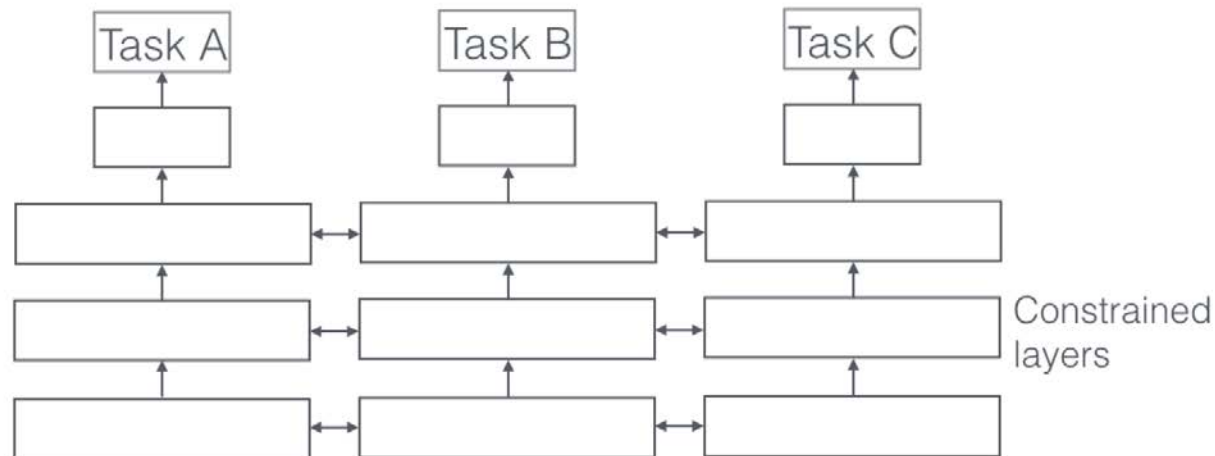
## ● Hard parameter sharing

- Greatly reduce the risk of overfitting
- Similar concept of bagging



## ● Soft parameter sharing

- Take a role of regularization



# Learning task relationship with Regularization

## Notation

- Task  $T$ , for each task  $t$ , we have a model  $m_t$  with parameters  $a_t$  of dimensionality  $d$
- The parameter vector  $a_t$  and parameter matrix  $A$  is:

$$a_t = [a_{1,t}, \dots, a_{d,t}]$$
$$A = \begin{bmatrix} \vdots & \vdots & \vdots \\ a_{.1} & \cdots & a_{.T} \\ \vdots & \vdots & \vdots \end{bmatrix}$$

i-th features of the model for every task

Parameter  $a_j$  corresponding to the j-th model

## Learning task relationship

$$\Omega = \|\bar{a}\|^2 + \frac{\lambda}{T} \sum_{t=1}^T \|a_{.,t} - \bar{a}\|^2 \quad \text{where, } \bar{a} = (\sum_{t=1}^T a_{.,t})/T$$

$$\tilde{J}(\theta; X, y) = J(\theta; X, y) + \alpha \Omega(\theta)$$

Regularized objective function

Original objective function

Penalty term

# Parameter Sharing

- Parameter sharing forces sets of parameters to be equal
- Only a subset of parameters (the unique set) need to be stored in memory (memory efficient than parameter tying, especially in CNN)
- Example case : parameter sharing in a convolution layer of CNN

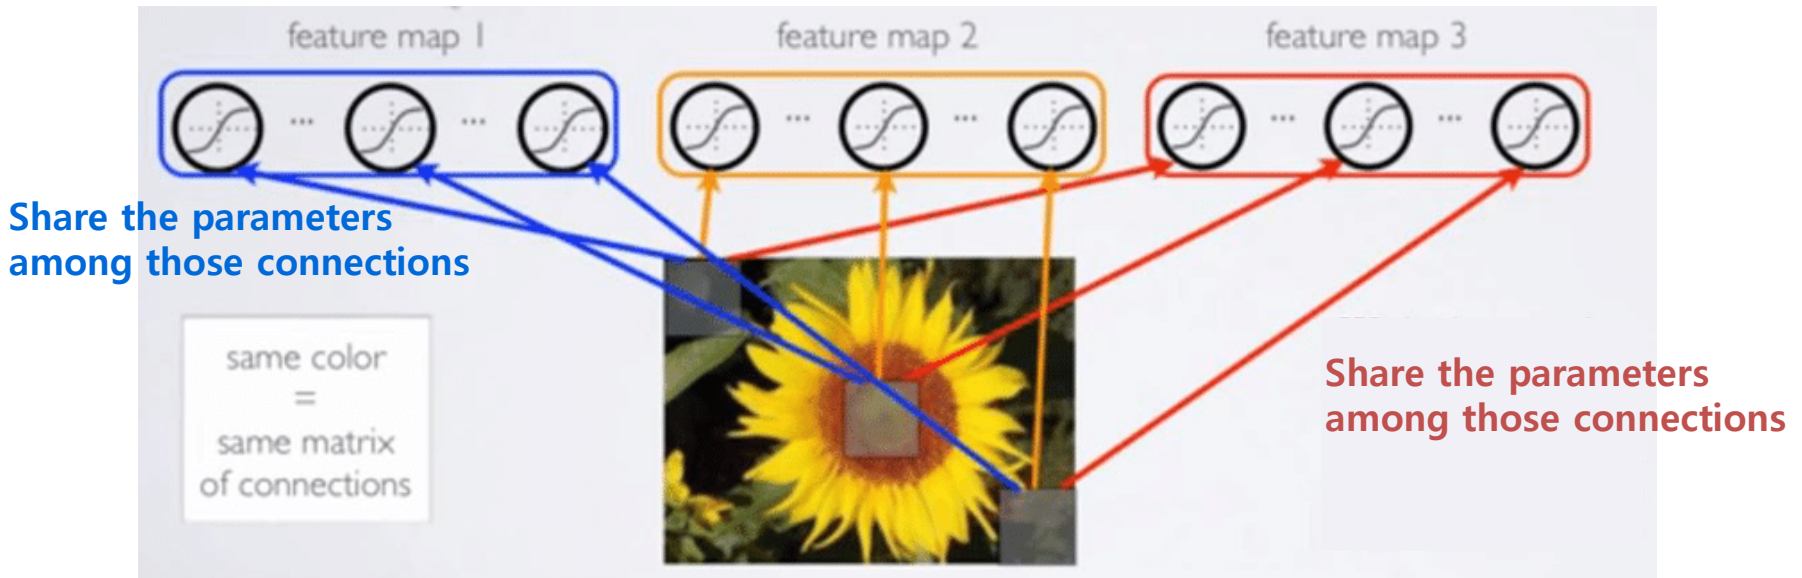



Image from Hugo Larochelle's lecture on YouTube (<https://www.youtube.com/watch?v=aAT1t9pZShM>)

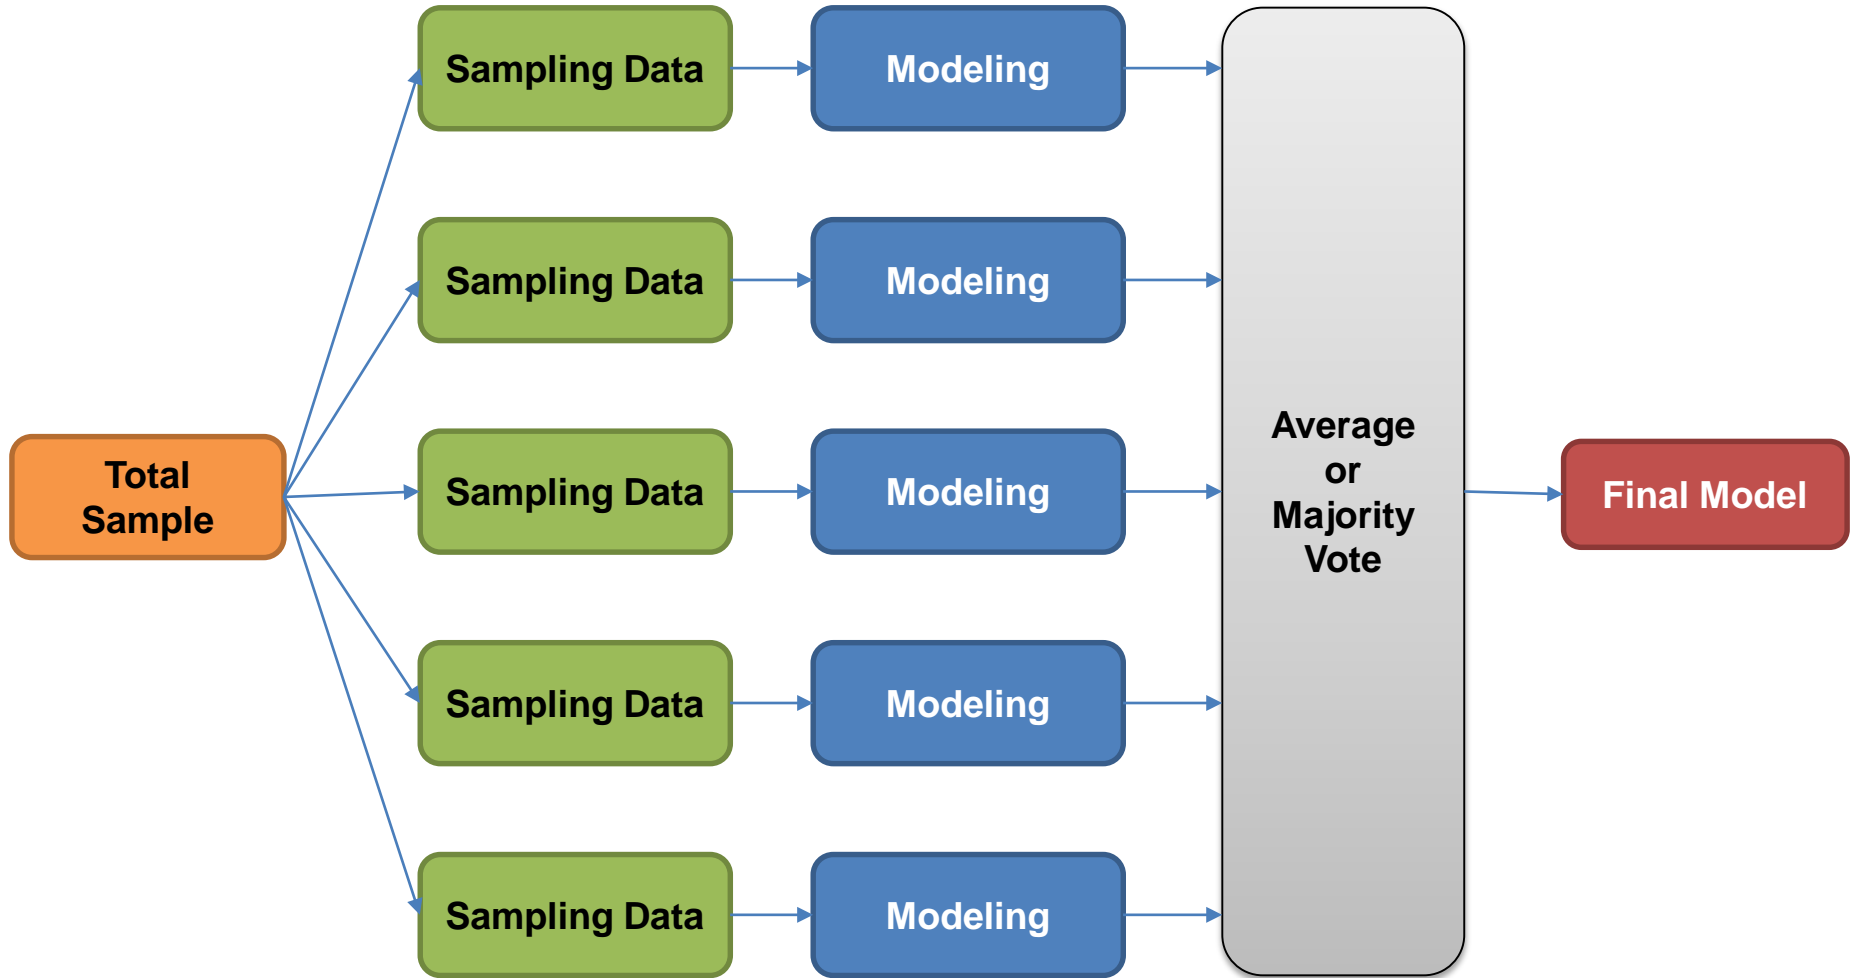
# Example of CNN filters

- 96 filters from AlexNet




# ***Bagging and Other Ensemble Methods***

# The Concept of Bagging







# Tree Based Bagging




<Model with Full dataset>




<Model No. 1>




<Model No. 2>



<Model No. 3>



<Model No. 4>



<Model No. 5>

# More Details of Bagging

- Short for *Bootstrap aggregation*
- Technique for reducing generalization error by combining several models
- a.k.a., *model averaging* or *ensemble methods*
- Procedure
  - Split the input data to  $K$  clusters with  $N'$  examples (sampling with replacement)
  - Train a classifier with a random sampled cluster
  - For testing, take each examples of test data to all classifier
  - Each classifier votes on the output, take majority

# Why dose the Bagging work? (1)

- Consider the  $K$  regression models (with minimize MSE)
- Suppose that each model make an error  $\epsilon_i$  on each model
  - Errors drawn from a zero-mean multivariate normal dist.

$$\text{variance } v = \mathbb{E}[\epsilon_i^2] \quad \text{covariance } c = \mathbb{E}[\epsilon_i \epsilon_j]$$

- Error made by the average prediction of models is:  $\frac{1}{k} \sum_i \epsilon_i$
- The expected squared error of the ensemble predictor is:

$$\begin{aligned} \mathbb{E} \left[ \left( \frac{1}{k} \sum_i \epsilon_i \right)^2 \right] &= \frac{1}{k^2} \mathbb{E} \left[ \sum_i \left( \epsilon_i^2 + \sum_{i \neq j} \epsilon_i \epsilon_j \right) \right] \\ &= \frac{1}{k^2} \{ k \mathbb{E}[\epsilon_i^2] + k(k-1) \mathbb{E}[\epsilon_i \epsilon_j] \} \\ &= \frac{1}{k} v + \frac{k-1}{k} c \end{aligned}$$

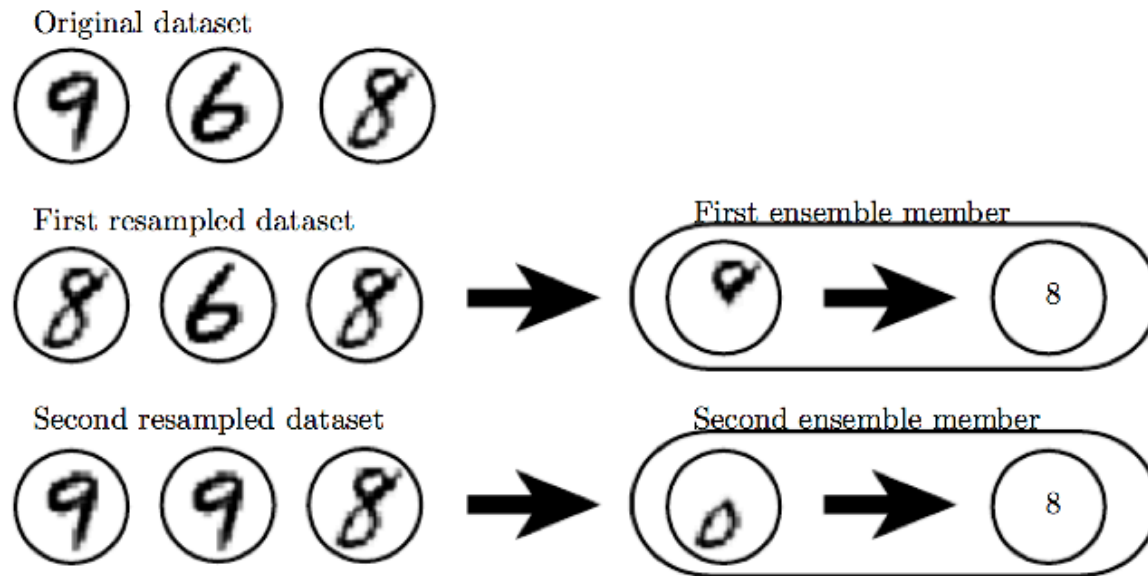
# Why dose the Bagging work? (2)

- The expected squared error of the ensemble predictor is:

$$\frac{1}{k}v + \frac{k-1}{k}c$$

- If the errors are perfectly correlates,  $c = v$ , it will not work at all
- If the errors are perfectly uncorrelated,  $c = 0$ , error will be only  $\frac{1}{k}v$


- Intuitive example



# Why dose the Bagging work? (3)

## ● Machine Learning Example

- Trend regression on the data  
Ozone-Temperature
- Gray line is regression line with  
each samples
- Red line is average line



## ● The case should not apply bagging

- When the sample data size is small
- When the data is noisy
- When the data have dependencies

# Tacit rules in Bagging

- OOB (Out-of-Bag) sampling

- Special rule for sampling with replacement
- If we sample the example with random sampling replacement, the selecting probability of each example is:

$$1 - \left(1 - \frac{1}{N}\right)^N \quad \text{If } N \text{ is large enough, then } \lim_{n \rightarrow \infty} \left\{1 - \left(1 - \frac{1}{N}\right)^N\right\} = 1 - \frac{1}{e} \approx 0.632$$

- Bagging in Neural Networks

- Random initialization
- Random selection of minibatches
- Differences in hyperparameter
- ...

- **Usually discouraged when benchmarking algorithms for scientific papers, because of its power and reliability**

- It's the benefit from the price of increased computations and memory



# Dropout

## References:

- [1] Nitish Srivastava et al., "Dropout: A Simple Way to Prevent Neural Networks from Overfitting", *Journal of Machine Learning Research* 15 (2014), 1929-1958
- [2] Hinton, Geoffrey E., et al., "Improving neural networks by preventing co-adaptation of feature detectors." *arXiv preprint arXiv:1207.0580* (2012).
- [3] Krizhevsky, Alex et al., "Imagenet classification with deep convolutional neural networks." *Advances in neural information processing systems*. 2012.
- [4] Wan, Li, et al. "Regularization of neural networks using dropconnect." *Proceedings of the 30th international conference on machine learning (ICML-13)*. 2013.
- [5] Baldi, Pierre, and Peter Sadowski. "The dropout learning algorithm." *Artificial intelligence* 210 (2014): 78-122.

# Overfitting

- Excessive focus on train data, resulting in worse results on actual test data





# Solutions for Overfitting

- **Regularization**

- L1-norm penalty
- L2-norm penalty

- **Data augmentation**

- **Dropout (2012)**


- A method of bagging applied to neural networks
- An inexpensive but powerful method of regularizing a broad family of models

- **Batch Normalization (2015)**

- It is described further in Chapter 8. Optimization

# Research in Dropout



- First proposed by G.E. Hinton (2012)<sup>[2]</sup>
- Became popular by *AlexNet* (2012)<sup>[3]</sup>
  - Winner in ILSVRC-2012 (ImageNet challenge)
  - AlexNet outperforms the other models at most 2x
  - CNN model with ReLU, Dropout, Data augmentation, GPU
  - Applied the dropout at Full-Connected layer




- Reinforced by S. Nitish (2014)<sup>[1]</sup>
  - Strengthen the theoretical background, extend to convolutional layer

# Dropout


- A technique that omits a portion of the network
  - We can surely improve the performance by model combination like as Bagging concept
  - However, if neural network is too deep to build the multiple models, it might be costly and inefficient
  - Also, it takes long time to inference the input with multiple models
- Dropout addressed the two problems
  - Omit the neurons, to mimic the voting in ensemble technique, instead of building the multiple models
  - Product the probability that a neurons will survive to weight, at inference level



# Dropout



(a) Standard Neural Net




(b) After applying dropout.

# Effect of Dropout

- Avoid the co-adaptation[2]

- Co-adaptation: the trend that some neurons tend to represent similar features
- Capture the clear features by avoiding co-adaptation




(b) After applying dropout.




# Effect of Dropout

- Effects on image classification models for MNIST




(a) Without dropout




(b) Dropout with  $p = 0.5$ .

# Dropout Modeling



(a) Standard network





(b) Dropout network

$$\begin{aligned} z_i^{(l+1)} &= \mathbf{w}_i^{(l+1)} \mathbf{y}^l + b_i^{(l+1)}, \\ y_i^{(l+1)} &= f(z_i^{(l+1)}), \end{aligned}$$

$$\begin{aligned} r_j^{(l)} &\sim \text{Bernoulli}(p), \\ \tilde{\mathbf{y}}^{(l)} &= \mathbf{r}^{(l)} * \mathbf{y}^{(l)}, \\ z_i^{(l+1)} &= \mathbf{w}_i^{(l+1)} \tilde{\mathbf{y}}^l + b_i^{(l+1)}, \\ y_i^{(l+1)} &= f(z_i^{(l+1)}). \end{aligned}$$



# Dropout Modeling

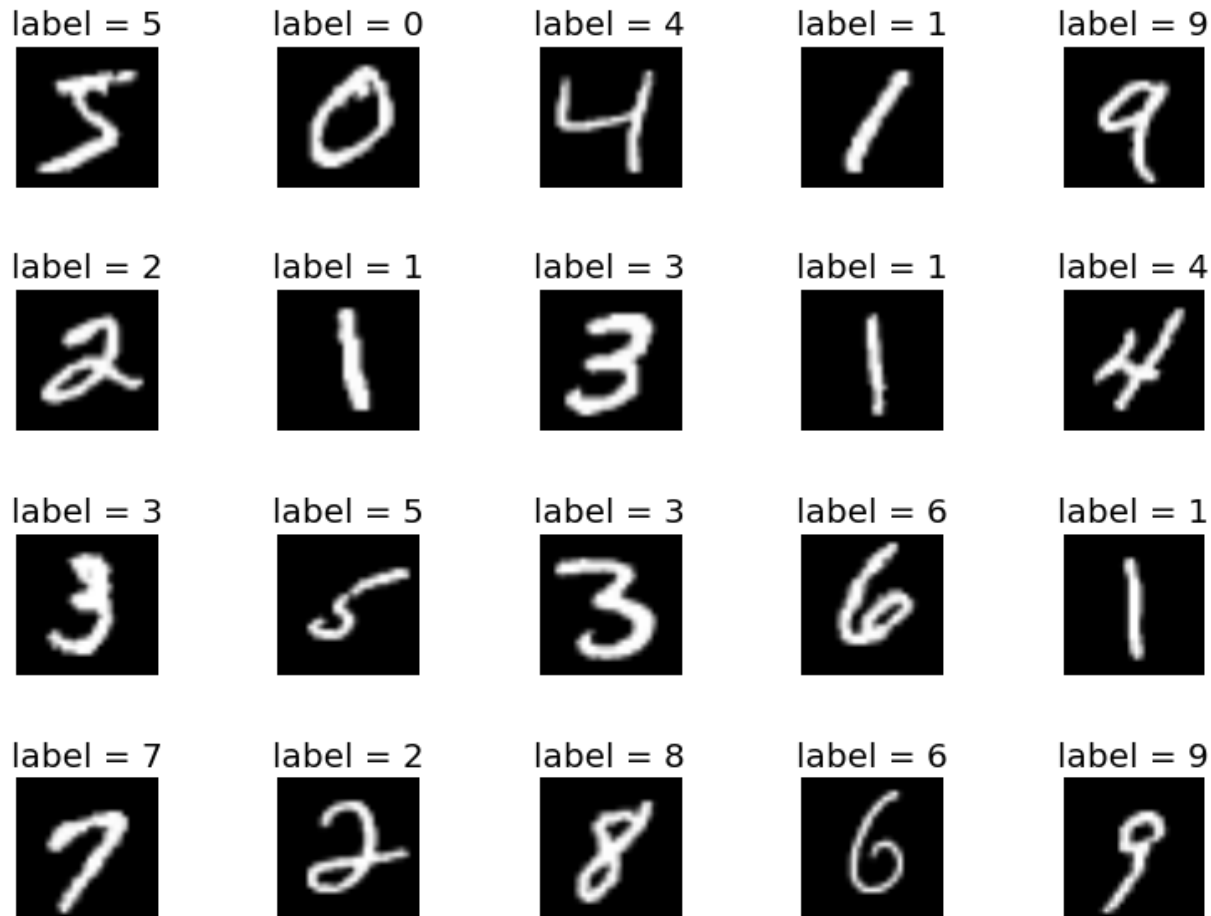


Ensemble of Sub-Networks



# Dropout Performance

- MNIST : a standard toy data set of handwritten digits



# Dropout Performance


## ● MNIST results

Method	Unit Type	Architecture	Error %
Standard Neural Net (Simard et al., 2003)	Logistic	2 layers, 800 units	1.60
SVM Gaussian kernel	NA	NA	1.40
Dropout NN	Logistic	3 layers, 1024 units	1.35
Dropout NN	ReLU	3 layers, 1024 units	1.25
Dropout NN + max-norm constraint	ReLU	3 layers, 1024 units	1.06
Dropout NN + max-norm constraint	ReLU	3 layers, 2048 units	1.04
Dropout NN + max-norm constraint	ReLU	2 layers, 4096 units	1.01
Dropout NN + max-norm constraint	ReLU	2 layers, 8192 units	0.95
Dropout NN + max-norm constraint (Goodfellow et al., 2013)	Maxout	2 layers, (5 × 240) units	0.94

# Dropout Performance


- MNIST results

- With fixed dropout rate, for different architectures




# Dropout another effect: Sparsity

- Can capture salient features
  - Makes neurons more sparse
  - Prevent co-adaptation




(a) Without dropout




(b) Dropout with  $p = 0.5$ .

# Hyper parameter $p$ : the Dropout rate

- (a) fixed number of neurons, variable  $p$ 
  - Relatively constant test error on 0.4~0.8  $\rightarrow$  usually use  $p = 0.5$
- (b) fixed value of  $pn$ 
  - Lower test error on low  $p \rightarrow$  increase number of neurons for low  $p$



(a) Keeping  $n$  fixed.




(b) Keeping  $pn$  fixed.

Figure 9: Effect of changing dropout rates on MNIST.

# Effect of data set size on Dropout


- Dropout is more powerful for larger dataset




# DropConnect

- A generalization of Dropout [4]


- Dropout omit the all connections which related to the unit
- DropConnect only omit the connections and leave the unit alive



Original



Dropout



DropConnect

# DropConnect Modeling

## ● Notation

- $h_i(I)$  is output value of a neuron
- $w_{ij}$  is weight of connections  $j$  to  $i$
- $I_j$  is input of a neuron

## ● The standard neural networks

$$h_i(I) = \sum_{j=1}^n w_{ij} I_j$$

## ● Dropout neural networks

$$h_i(I) = \sum_{j=1}^n w_{ij} r_j I_j$$

## ● DropConnect neural networks

$$h_i(I) = \sum_{j=1}^n r_{ij} w_{ij} I_j$$



# DropConnect Performance

neuron	model	error(%) 5 network	voting error(%)
<i>relu</i>	No-Drop	$1.62 \pm 0.037$	1.40
	Dropout	$1.28 \pm 0.040$	1.20
	DropConnect	$1.20 \pm 0.034$	<b>1.12</b>
<i>sigmoid</i>	No-Drop	$1.78 \pm 0.037$	1.74
	Dropout	$1.38 \pm 0.039$	<b>1.36</b>
	DropConnect	$1.55 \pm 0.046$	1.48
<i>tanh</i>	No-Drop	$1.65 \pm 0.026$	1.49
	Dropout	$1.58 \pm 0.053$	1.55
	DropConnect	$1.36 \pm 0.054$	<b>1.35</b>

# ***Adversarial Training***

References:

- [1] Goodfellow, Ian J., Jonathon Shlens, and Christian Szegedy. "Explaining and harnessing adversarial examples." *arXiv preprint arXiv:1412.6572* (2014).
- [2] Szegedy, Christian, et al. "Intriguing properties of neural networks." *arXiv preprint arXiv:1312.6199* (2013).

# Adversarial Examples

## Definition

- Applying an imperceptible non-random perturbation to a test image, it is possible to arbitrarily change the network's prediction [2]

## Examples

- With the same predict model,



Y = "panda"

With 0.577  
confidence



Y = "gibbon"

With 0.993  
confidence

# Adversarial Examples

## Definition

- Applying an imperceptible non-random perturbation to a test image, it is possible to arbitrarily change the network's prediction [2]

## Examples

- With the same predict model,



Y = "panda"

With 0.577  
confidence

× 0.07





=



Y = "gibbon"

With 0.993  
confidence

# Adversarial Examples



Correctly predicted sample

Adversarial examples

# Adversarial Training

- Formal description[2]

$$\begin{array}{llll} \text{Minimize } \|\eta\|_2 & \text{subject to} & f(x + \eta) = l & \text{noise : } \eta \\ & & x + \eta \in [0,1]^m & \text{label : } l \\ & & f(x) \neq l & \end{array}$$

- More general description for neural network training[1]

Training as a regularization term

$$\tilde{J}(\theta, x, y) = \alpha J(\theta, x, y) + (1 - \alpha) J(\theta, x + \eta)$$

$$\text{where, } \eta = \epsilon \text{sign}(\nabla_x J(\theta, x, y))$$



# Current research

Different lighting conditions:



all: right



all: right



all: right



all: 1



all: 3




all: 5



DRV\_C1: left



DRV\_C2: left



DRV\_C3: left



MNI\_C1: 8



MNI\_C2: 5



MNI\_C3: 7

# ***The next Deep Learning Seminar***

## **Chapter 8. Optimization for Training Deep Models**

- 8.1 How Learning Differs from Pure Optimization**
- 8.2 Challenges in Neural Network Optimization**
- 8.3 Basic Algorithms**
- 8.4 Parameter Initialization Strategies**
- 8.5 Algorithms with Adaptive Learning Rates**
- 8.6 Approximation Second-Order Methods**
- 8.7 Optimization Strategies and Meta-Algorithms**



