

- Part 2 -

Department of Information and Communication Engineering

2017.11.22

Contents

Chapter 10. Sequence Modeling

● Part 1

- 10.1 Unfolding Computational Graphs
- 10.2 Recurrent Neural Networks
- 10.3 Bidirectional RNNs
- 10.4 Encoder-Decoder Sequence-to-Sequence Architecture
- 10.5 Deep Recurrent Networks
- 10.6 Recursive Neural Networks

● Part 2

- 10.7 The challenge of Long-term dependencies
- 10.8 Echo State Networks
- 10.9 Leaky Units and Other strategies for Multiple Time Scales
- 10.10 The Long Short-Term Memory and Other Gated RNNs
- 10.11 Optimization for Long-Term Dependencies
- 10.12 Explicit Memory

Chapter 10. Sequence Modeling

● Part 1

- 10.1 Unfolding Computational Graphs
- 10.2 Recurrent Neural Networks
- 10.3 Bidirectional RNNs
- 10.4 Encoder-Decoder Sequence-to-Sequence Architecture
- 10.5 Deep Recurrent Networks
- 10.6 Recursive Neural Networks


● Part 2

- 10.7 The challenge of Long-term dependencies
- 10.8 Echo State Networks
- 10.9 Leaky Units and Other strategies for Multiple Time Scales
- 10.10 The Long Short-Term Memory and Other Gated RNNs
- 10.11 Optimization for Long-Term Dependencies
- 10.12 Explicit Memory

The Challenge of Long-term dependencies

Long-Term Dependencies

- RNN might be able to connect previous information to present task



- But, can they?




Image from colah's blog on github (<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>)

Repeatedly multiplying weight

● The simple case

- Let model's parameter matrix W
- In recurrent case, k step is equivalent to W^k
- Suppose that W has an eigen-decomposition


$$W = V \text{diag}(\lambda) V^{-1}$$

- $W^k = (V \text{diag}(\lambda) V^{-1})^k = V \text{diag}(\lambda)^k V^{-1}$
- Any eigenvalue λ_i that are not near an absolute value of 1 will either explode or vanishing

● In general models

- $h^t = W^T h^{(t-1)} = (W^t)^T h^{(0)} = Q^T \Lambda^t Q h^{(0)}$
- Long-term dependencies arises from the exponentially smaller weights given to long-term interactions compared to short-term ones.

Repeatedly multiplying weight




Echo State Network


Echo State Network

- Simple trick to make RNN much easier to learn
- Initialize the connections in the RNN **randomly** in such a way that it has big reservoir of coupled oscillators
- **Do not consider** input to hidden or hidden to hidden learn
- The only thing have learn is **how to couple the output to the oscillators (state to output)**
- In the end, it just fit in a linear model
- Similar with SVM
 - Reservoir can be considered as kernel
 - State to output connection is classifier

Echo State Network



Transforming sine wave example



Pros and Cons

● Pros

- ESN can be **trained very fast** because they just fit linear model
- We can do many experiments
- ESN can do impressive modeling of 1-dimensional time-series

● Cons


- ESN **need many more units in reservoir** for more complex problems
- Initialization is very important to performance

● Applications


- ESN can be used for initialization of other networks

Long-Short Term Memory

RNN Review



$$Y^t = [\text{---} \text{---} \text{---} \text{vector representation} \text{---} \text{---}]$$



$$X^t = [\text{---} \text{---} \text{---} \text{vector representation} \text{---} \text{---}]$$

RNN Review

$$Y^t = [\text{---} \text{---} \text{---} \text{vector representation} \text{---} \text{---}]$$

output layer (k)


$$\uparrow W^{out} = (w_{kj}^{out})$$

hidden layer (j)

$$\curvearrowright W = (w_{jj'})$$


$$\uparrow W^{in} = (w_{ji}^{in})$$

input layer (i)



$$X^t = [\text{---} \text{---} \text{---} \text{vector representation} \text{---} \text{---}]$$

RNN Review



RNN Review - Forward propagation

- input of hidden layer

$$u_j^t = \sum_i w_{ji}^{(in)} x_i^t + \sum_j w_{jj} z_j^{t-1}$$

- output of hidden layer





$$z_j^t = f(u_j^t) \quad z^t = f(W^{in} X^t + W z^{t-1})$$

- input of output layer

$$v_k^t = \sum_j w_{kj}^{(out)} z_j^t$$

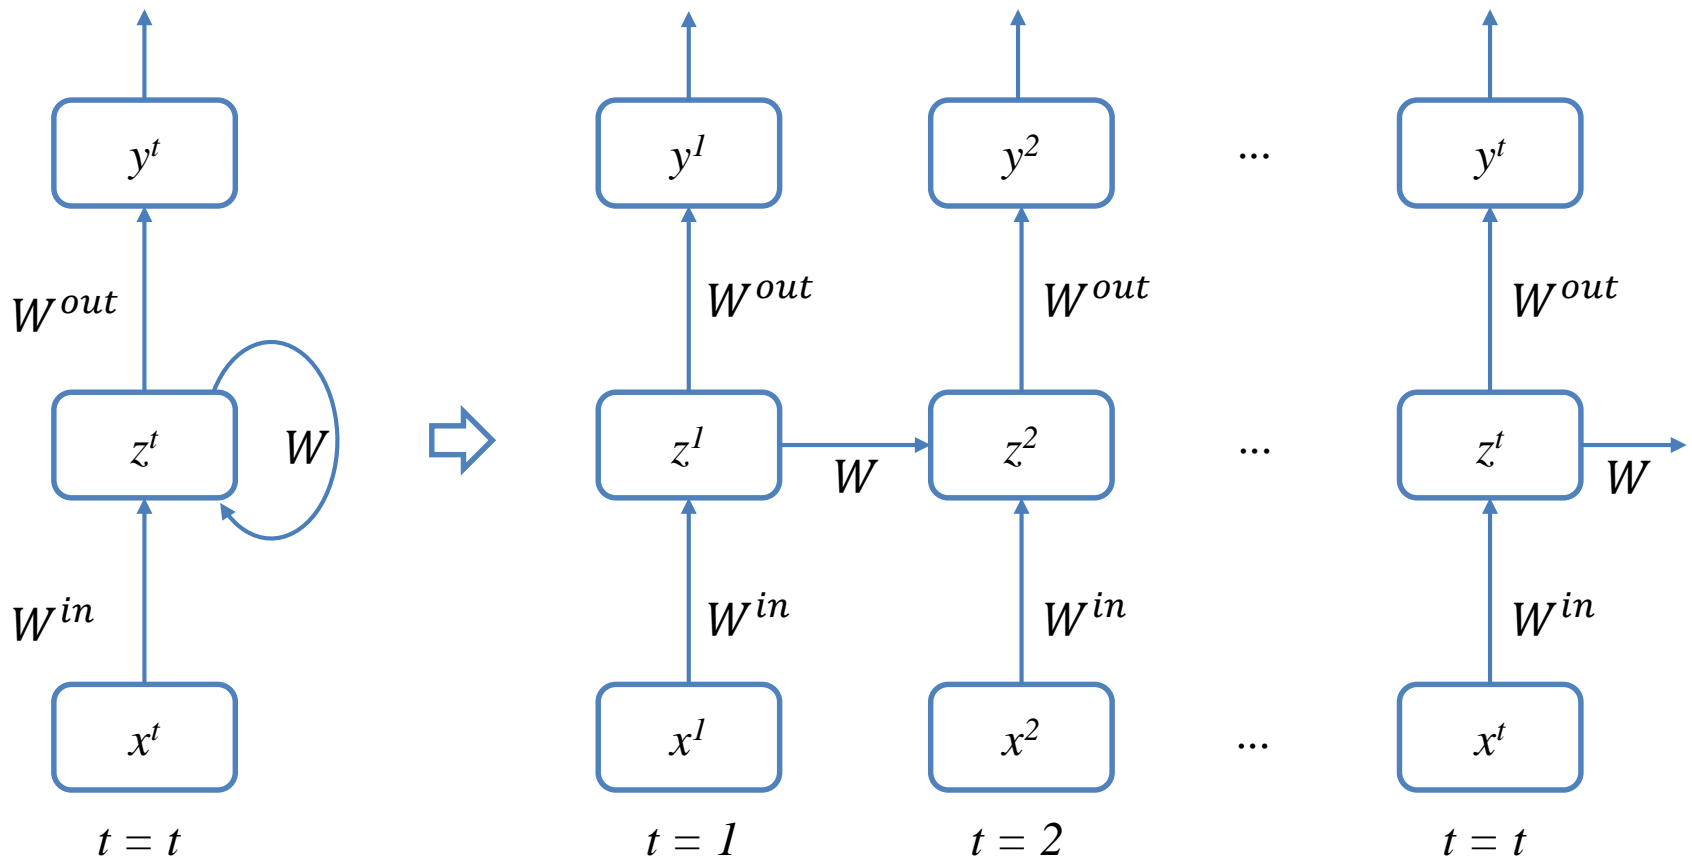
- output of output layer

$$y^t = f^{out}(v^t) = f^{out}(W^{out} z^t)$$



Back Propagation Through Time(BPTT) (1)

- Faster and more simple than RTRL(RealTime Recurrent Learning)
- Deploy the RNN with time direction (like Feedforward NN)
 - assumes that each layer is separated by time

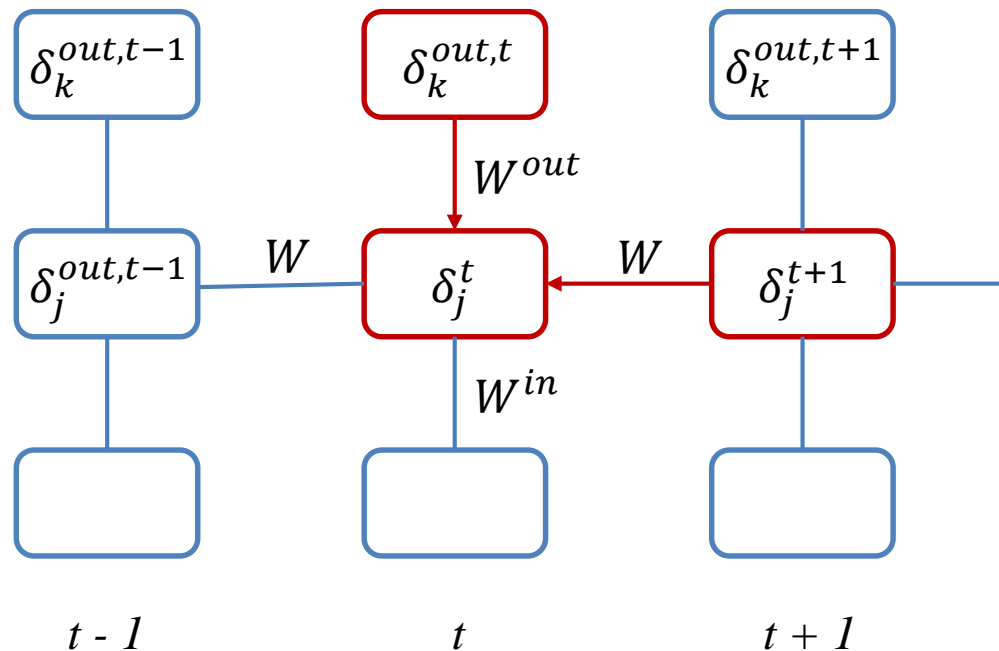


Back Propagation Through Time(BPTT) (2)

- In Feedforward NN, we can find *delta* at layer *l* as :

$$\delta_j^l \equiv \frac{\partial E}{\partial u_j^l} \quad \delta_j^l = \sum_k w_{kj}^{l+1} \delta_k^{l+1} f'(u_j^l)$$


- In RNN, we *delta* is affected by :



LSTM (Long Short-Term Memory)

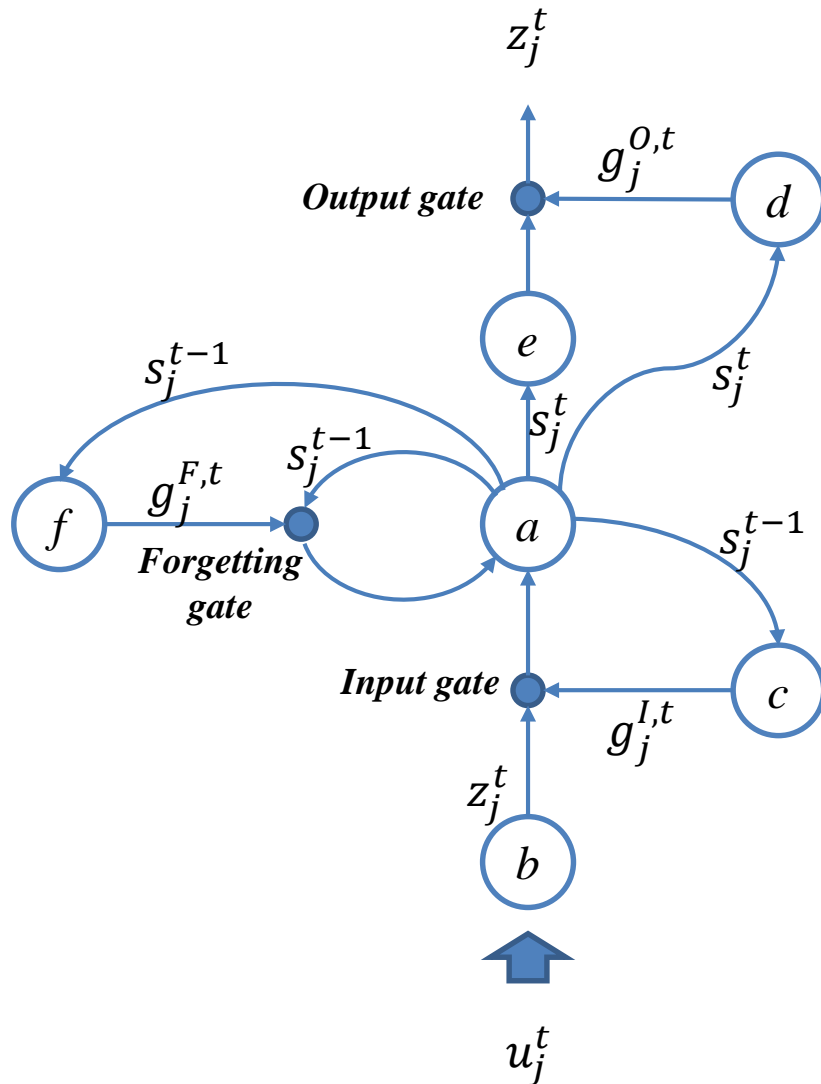
- Designed to overcome RNN's limitation

- in practice, pure RNN dose not contain the information of whole inputs
- pure RNN memorize the data up to $(t-10)$



- a. memory cell
- b. a RNN's node
- c. node for input gate
- d. node for output gate
- e. node for activation function of memory cell output
- f. node for forgetting gate

LSTM gate operation



output gate :

$$g_j^{O,t} \in [0,1]$$

$$g_j^{O,t} \cdot s_j^t$$

forgetting gate :

$$g_j^{F,t} \in [0,1]$$

$$g_j^{F,t} \cdot s_j^{t-1}$$

input gate :

$$g_j^{I,t} \in [0,1]$$

$$g_j^{I,t} \cdot z_j^t$$

LSTM forward propagation (1)

- before node b ,

$$u_j^t = \sum_i w_{ji}^{(in)} x_i^t + \sum_j w_{jj} z_j^{t-1}$$

- from node f ,





$$g_j^{F,t} = f_f(u_j^{F,t}) = f_f\left(\sum_i w_{ji}^{F,in} x_i^t + \sum_{j'} w_{jj'}^{F,z} z_{j'}^{t-1} + w_j^F s_j^{t-1}\right)$$

- from node c ,

$$g_j^{I,t} = f_c(u_j^{I,t}) = f_c\left(\sum_i w_{ji}^{I,in} x_i^t + \sum_{j'} w_{jj'}^{I,z} z_{j'}^{t-1} + w_j^I s_j^{t-1}\right)$$

- so, memory cell a is:


$$s_j^t = g_j^{F,t} s_j^{t-1} + g_j^{I,t} f_b(u_j^t)$$



LSTM forward propagation (2)


- from node d ,

$$g_j^{0,t} = f_d(u_j^{0,t}) = f_d\left(\sum_i w_{ji}^{0,in} x_i^t + \sum_{j'} w_{jj'}^0 z_{j'}^{t-1} + w_j^0 s_j^t\right)$$



- so the output z_j^t is :

$$z_j^t = g_j^{0,t} f_e(s_j^t)$$



- usually, g is activation function with *Logistic Sigmoid Func.*

LSTM Back propagation (1)

- In general feedforward NN structure,

$$\delta_j^l = \sum_k \delta_k^{l+1} \frac{\partial u_k^{l+1}}{\partial u_j^l}$$

- In LSTM, z_j^t is the only value which affected by external node

$$v_k^t = \sum_j w_{kj}^{out} z_j^t$$

- Since z_j^t is as follows,

$$z_j^t = g_i^{0,t} f_e(s_j^t)$$


- so, gradient of v_k^t is:

$$\frac{\partial v_k^t}{\partial u_j^{0,t}} = w_{kj}^{out} f'(u_j^{0,t}) f_e(s_j^t)$$

LSTM Back propagation (2)

- It means, gradient of temporal node g is :


$$\frac{\partial v_k^t}{\partial u_j^{o,t}} = w_{kj}^{out} f'(u_j^{o,t}) f_e(s_j^t)$$



- At node d , gradient is :


$$\varepsilon_j^t = \sum_k w_{kj}^{out} \delta_k^{out,t} + \sum_{j'} w_{jj'} \delta_{j'}^{t+1}$$

$$\delta_j^{o,t} = f'(u_j^{o,t}) f_d(s_j^t) \varepsilon_j^t$$



- At node e , gradient is :


$$\widetilde{\delta}_j^t = g_j^{o,t} f'_d(s_j^t) \varepsilon_j^t$$



LSTM Back propagation (3)


- At node a , it has 5 connection, and gradient is :

$$\delta_j^{cell,t} = \widetilde{\delta}_j^t + g_j^{F,t-1} \delta_j^{cell,t+1} + w_j^F \delta_j^{F,t+1} + w_j^I \delta_j^{I,t+1} + w_j^O \delta_j^{O,t}$$




- At node b , gradient is :

$$\delta_j^t = g_j^{I,y} f'_c(u_j^t) \delta_j^{cell,t}$$




- At node f , gradient is :

$$\delta_j^{F,t} = f'(u_j^{F,t}) s_j^{t-1} \delta_j^{cell,t}$$




- At node c , gradient is :


$$\delta_j^{I,t} = f'(u_j^{I,t}) f(u_j^t) \delta_j^{cell,t}$$



Other Method: GRU (Gated Recurrent Unit)



LSTM Gating



GRU Gating

- More simpler version than LSTM

- Use only 2 gate

- Both method have pros and cons

Image form, Chung, Junyoung, et al. "Empirical evaluation of gated recurrent neural networks on sequence modeling." (2014)

Optimization for Long-Term Dependencies

Two major issues in LTD optimizations

- Gradient Exploding


- Deal with Gradient clipping

- Gradient Vanishing

- LSTM


Cliffs and Exploding Gradient

- Cliff is also a common issue in such as those computed by a recurrent net over many time steps
 - resulting from the multiplication of several parameters
- The gradient update step can move the parameter extremely far (exploding gradient)
 - losing most of optimization works that has been done




Gradient Clipping for Handling Cliffs

$$\Delta\theta_i \leftarrow \Delta\theta_i \frac{\Delta\text{MAX}}{\max(|\Delta\theta_i|, \Delta\text{MAX})}$$




RNN with LSTM



Explicit Memory

Explicit Memory

- Many researches using memory are ongoing to address LTD
- They learned mechanism of control the memory, and deciding where to read from and where to write to



Memory Networks

● You can find the paper at:

- <https://arxiv.org/abs/1410.3916>

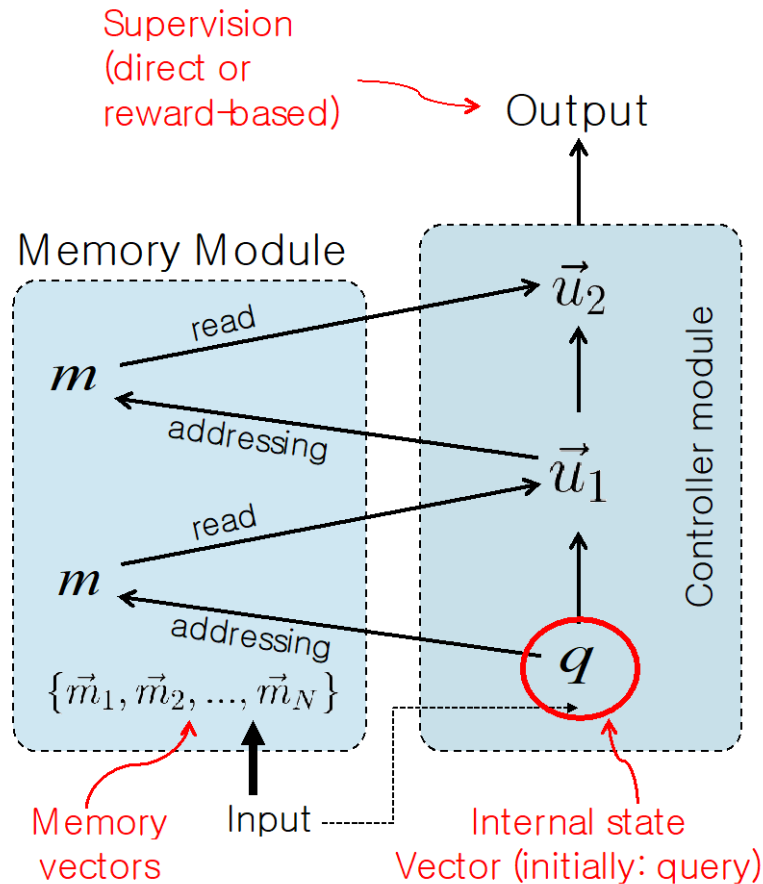


Image from Jason Weston, "Memory network tutorial", at ICML 2016

Neural Turing Machine

● You can find the paper at:

- <https://arxiv.org/abs/1410.5401>


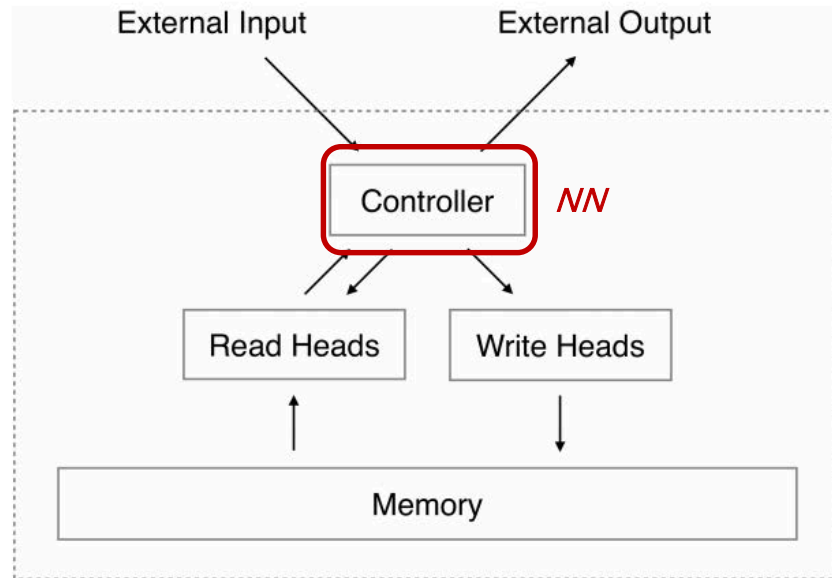


Image from https://norman3.github.io/papers/docs/neural_turing_machine.html

The next Deep Learning Seminar

Chapter 11. Practical Methodology

- 11.1 Performance Metrics**
- 11.2 Default Baseline Models**
- 11.3 Determining Whether to Gather More Data**
- 11.4 Selecting Hyperparameters**
- 11.5 Debugging Strategies**
- 11.6 Examples: Multi-Digit Number Recognition**

Chapter 12. Applications

- 12.1 Large Scale Deep Learning**
- 12.2 Computer Vision**
- 12.3 Speech Recognition**
- 12.4 Natural Language Processing**
- 12.5 Other Applications**

